

Road Traffic Dataset

Total lines of code 102.

FINISHED

Data loading and running of some SQL queries took between 1 and 2 minutes as the data set consists of more than 20 million rows.

Took 3 sec. Last updated by anonymous at March 31 2017, 1:26:15 PM.

```
1 %pyspark
2 from pandas import Series, DataFrame
3 import pandas as pd
4 import numpy as np
5 import glob,os
6 path = "/Users/Kapil/Downloads/traffic_feb_june"
7 roadtraffic = pd.concat(map(pd.read_csv, glob.glob(os.path.join(path, "*.csv"))))
```

FINISHED

Took 52 sec. Last updated by anonymous at April 16 2017, 9:44:29 PM.

```
1 %pyspark
2 inputPath = "/Users/Kapil/Downloads/traffic_feb_june"
3 roadtraffic2 = sqlContext.read.format("com.databricks.spark.csv").option("header", "true")
```

FINISHED

Took 2 min 0 sec. Last updated by anonymous at April 16 2017, 9:46:43 PM.

```
1 %pyspark
2 roadtraffic2.registerTempTable("road_traffic")
```

FINISHED

Took 0 sec. Last updated by anonymous at April 16 2017, 10:29:53 PM.

```
1 %pyspark
2 print(roadtraffic.count())
```

FINISHED

```
status                20713165
avgMeasuredTime        20713165
avgSpeed               20713165
extID                  20713165
medianMeasuredTime     20713165
TIMESTAMP              20713165
vehicleCount           20713165
_id                    20713165
REPORT_ID              20713165
dtype: int64
```

Took 11 sec. Last updated by anonymous at March 31 2017, 12:21:50 PM.

```
1 %pyspark
2 roadtraffic[-5:]
```

FINISHED

	status	avgMeasuredTime	avgSpeed	extID	medianMeasuredTime	\
16938	OK	112	36	623	112	
16939	OK	112	36	623	112	
16940	OK	112	36	623	112	
16941	OK	112	36	623	112	
16942	OK	112	36	623	112	

	TIMESTAMP	vehicleCount	_id	REPORT_ID
16938	2014-09-30T23:35:00	0	28062086	210199
16939	2014-09-30T23:40:00	0	28062468	210199
16940	2014-09-30T23:45:00	0	28062917	210199
16941	2014-09-30T23:50:00	0	28063308	210199
16942	2014-09-30T23:55:00	0	28063757	210199

Took 0 sec. Last updated by anonymous at March 31 2017, 12:21:55 PM. (outdated)

```

1 %pyspark
2 import re
3 roadtraffic['hour'] = roadtraffic['TIMESTAMP'].str[11:13]
4 roadtraffic['minutes'] = roadtraffic['TIMESTAMP'].str[14:16]
5 roadtraffic['date'] = roadtraffic['TIMESTAMP'].str[0:10]
6 roadtraffic['months'] = roadtraffic['TIMESTAMP'].str[5:7]

```

FINISHED

Took 39 sec. Last updated by anonymous at March 31 2017, 12:36:24 PM.

```

1 %pyspark
2 roadtraffic[-5:]

```

FINISHED

	status	avgMeasuredTime	avgSpeed	extID	medianMeasuredTime	\
16938	OK	112	36	623	112	
16939	OK	112	36	623	112	
16940	OK	112	36	623	112	
16941	OK	112	36	623	112	
16942	OK	112	36	623	112	

	TIMESTAMP	vehicleCount	_id	REPORT_ID	hour	minutes	\
16938	2014-09-30T23:35:00	0	28062086	210199	23	35	
16939	2014-09-30T23:40:00	0	28062468	210199	23	40	
16940	2014-09-30T23:45:00	0	28062917	210199	23	45	
16941	2014-09-30T23:50:00	0	28063308	210199	23	50	
16942	2014-09-30T23:55:00	0	28063757	210199	23	55	

	date	months
16938	2014-09-30	09
16939	2014-09-30	09
16940	2014-09-30	09
16941	2014-09-30	09
16942	2014-09-30	09

Took 0 sec. Last updated by anonymous at March 31 2017, 12:36:32 PM.

```

1 %pyspark
2 roadtraffic.info()

```

FINISHED

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 20713165 entries, 0 to 16942
Data columns (total 13 columns):
status                object
avgMeasuredTime       int64
avgSpeed              int64
EXTID                 int64

```

```

EXTID          int64
medianMeasuredTime  int64
TIMESTAMP      object
vehicleCount    int64
_id            int64
REPORT_ID       int64
hour           object
minutes        object
date           object
months         object
dtypes: int64(7), object(6)
memory usage: 2.2+ GB

```

Took 8 sec. Last updated by anonymous at March 31 2017, 12:36:45 PM.

```

1 %pyspark FINISHED
2 def get_stats(group): return {'min': group.min(), 'max': group.max(), 'count': group.coun
  group.mean()}
3 grouped_avgspeed_byhour = roadtraffic['avgSpeed'].groupby(roadtraffic['hour'])
4 grouped_avgspeed_byhour.apply(get_stats).unstack()

```

	count	max	mean	min
hour				
00	781546.0	149.0	48.058822	0.0
01	887090.0	149.0	48.348357	0.0
02	904458.0	149.0	48.544265	0.0
03	900229.0	150.0	48.353697	0.0
04	865603.0	150.0	46.639626	0.0
05	846795.0	150.0	42.866610	0.0
06	867806.0	149.0	40.894512	0.0
07	894057.0	150.0	42.059450	0.0
08	895654.0	150.0	42.058984	0.0
09	885253.0	150.0	41.679704	0.0
10	878261.0	149.0	41.382687	0.0
11	894081.0	149.0	41.330556	0.0
12	896679.0	149.0	40.876565	0.0
13	898122.0	149.0	40.752500	0.0

Took 6 sec. Last updated by anonymous at March 31 2017, 12:40:34 PM. (outdated)

```

1 %pyspark FINISHED
2 grouped_vehicleCount_byhour = roadtraffic['vehicleCount'].groupby(roadtraffic['hour'])
3 grouped_vehicleCount_byhour.apply(get_stats).unstack()

```

06	867806.0	121.0	5.720163	0.0
07	894057.0	99.0	5.069171	0.0
08	895654.0	79.0	5.036747	0.0
09	885253.0	65.0	5.164977	0.0
10	878261.0	67.0	5.303841	0.0
11	894081.0	77.0	5.357505	0.0
12	896679.0	90.0	5.627419	0.0
13	898122.0	97.0	6.022087	0.0
14	895562.0	94.0	6.050365	0.0
15	910856.0	85.0	4.940116	0.0
16	908012.0	74.0	3.505207	0.0
17	912939.0	71.0	2.405558	0.0
18	920914.0	68.0	1.797495	0.0
19	914818.0	78.0	1.529196	0.0
20	750271.0	86.0	1.252160	0.0

```

20  759571.0  80.0  1.232109  0.0
21  702923.0  85.0  0.782027  0.0
22  765726.0  58.0  0.398256  0.0
23  776410.0  67.0  0.749556  0.0

```

Took 4 sec. Last updated by anonymous at March 31 2017, 12:40:30 PM.

```

1 %pyspark FINISHED
2 grouped_avgMeasuredTime = roadtraffic['avgMeasuredTime'].groupby(roadtraffic['hour'])
3 grouped_avgMeasuredTime.apply(get_stats).unstack()

```

	count	max	mean	min
hour				
00	781546.0	3595.0	95.816972	0.0
01	887090.0	3587.0	95.751300	0.0
02	904458.0	3587.0	95.795096	0.0
03	900229.0	3587.0	96.949256	0.0
04	865603.0	3587.0	102.146635	0.0
05	846795.0	3587.0	114.061216	0.0
06	867000.0	3587.0	121.020027	0.0

Took 4 sec. Last updated by anonymous at March 31 2017, 12:39:16 PM. (outdated)

```

1 %pyspark FINISHED
2 grouped_avgspeed_bymonth = roadtraffic['avgSpeed'].groupby(roadtraffic['months'])
3 grouped_avgspeed_bymonth.apply(get_stats).unstack()

```

	count	max	mean	min
months				
02	1910192.0	149.0	42.935319	0.0
03	3485620.0	150.0	43.568671	0.0
04	3705591.0	150.0	43.705196	0.0
05	3681921.0	150.0	44.117599	0.0
06	793808.0	149.0	43.076908	0.0
08	3608396.0	150.0	44.692306	0.0
09	3527637.0	150.0	44.316682	0.0

Took 3 sec. Last updated by anonymous at March 31 2017, 12:43:16 PM. (outdated)

```

1 %pyspark FINISHED
2 grouped_vehicle_bymonth = roadtraffic['vehicleCount'].groupby(roadtraffic['months'])
3 grouped_vehicle_bymonth.apply(get_stats).unstack()

```

	count	max	mean	min
months				
02	1910192.0	111.0	3.380066	0.0
03	3485620.0	97.0	3.443799	0.0
04	3705591.0	111.0	2.854224	0.0
05	3681921.0	100.0	3.252956	0.0
06	793808.0	121.0	2.805493	0.0
08	3608396.0	107.0	3.151881	0.0
09	3527637.0	108.0	3.200058	0.0

Took 4 sec. Last updated by anonymous at March 31 2017, 12:45:06 PM. (outdated)

```

1 %pyspark FINISHED
2 grouped_avgTime_bymonth = roadtraffic['avgMeasuredTime'].groupby(roadtraffic['months'])
3 grouped_avgTime_bymonth.apply(get_stats).unstack()

```

	count	max	mean	min
months				
02	1910192.0	3587.0	103.188559	0.0
03	3485620.0	3648.0	104.820731	0.0
04	3705591.0	3595.0	105.554581	0.0
05	3681921.0	3656.0	109.032780	0.0
06	793808.0	3456.0	104.187155	0.0
08	3608396.0	3585.0	107.636983	0.0
09	3527637.0	3572.0	109.661455	0.0

Took 4 sec. Last updated by anonymous at March 31 2017, 12:45:28 PM. (outdated)

```
1 %sql
2 select * from road_traffic limit 15 --see the first 15 rows in the table
```

FINISHED



status	avgMeasuredTime	avgSpeed	extID	medianM
OK	141	93	672	141
OK	146	90	672	146
OK	155	84	672	155
OK	149	88	672	149
OK	146	90	672	146
OK	148	88	672	148
OK	150	87	672	150
OK	155	84	672	155

Took 1 sec. Last updated by anonymous at March 31 2017, 12:55:06 PM.

```
1 %sql
2 select count(_id) from road_traffic --count total number of rows
   in the table
```

FINISHED



count(_id)



20713165

Took 18 sec. Last updated by anonymous at March 31 2017, 12:50:56 PM. (outdated)

```
1 %sql
2 select distinct status from road_traffic--to see what distinct
   values are in status column
```

FINISHED



status



OK

Took 19 sec. Last updated by anonymous at March 31 2017, 12:53:08 PM. (outdated)

```
1 %sql
2 select distinct extID
3 from road_traffic
4 order by extID --to see what distinct values are in extID column
```

READY



eID



610

611

612

613

614

```
1 %sql
2 select count(eID) from (select distinct extID as eID from road_traffic group by extID) --
   number of distinct values in extID column
```

FINISHED




count(eID)

449

Took 18 sec. Last updated by anonymous at March 31 2017, 12:54:23 PM.

```
1 %sql
2 select avg(avgSpeed), hour(timestamp) as hour from road_traffic
3 group by hour(timestamp)
4 order by hour(timestamp)
```

FINISHED



Took 1 min 35 sec. Last updated by anonymous at March 31 2017, 1:03:43 PM. (outdated)

```
%sql
select avg(vehicleCount), hour(timestamp) as hour from road_traffic
group by hour(timestamp)
order by hour(timestamp)
```

FINISHED



All fields:

avg(vehicleCount) hour


Keys

hour ✕

Groups

Values


avg(vehicleCount) SUM ✕



Took 1 min 41 sec. Last updated by anonymous at April 16 2017, 10:36:47 PM. (outdated)

```
1 %sql
2 select avgSpeed, vehicleCount
3 from road_traffic
4
```

FINISHED



All fields:

avgSpeed vehicleCount

xAxis


avgSpeed ✕

yAxis

vehicleCount ✕

group

size ⓘ



Results are limited by 1000.


Took 0 sec. Last updated by anonymous at March 31 2017, 1:01:10 PM. (outdated)

```
1 %sql
2 select avg(vehiclecount), hour(timestamp) as hour
3 from road_traffic
4 group by hour(timestamp)
5 order by hour(timestamp)
```

FINISHED



settings ▼



Took 2 min 59 sec. Last updated by anonymous at March 31 2017, 1:05:12 PM. (outdated)


```
1 %sql
2 select date(timestamp) date, avg(avgSpeed)
3 from road_traffic
4 group by date(timestamp)
```

FINISHED

5 order by date(timestamp)



settings ▼



Took 1 min 30 sec. Last updated by anonymous at March 31 2017, 1:11:35 PM. (outdated)

```
1 %sql
2 select *
3 from road_traffic
4 order by road_traffic.timestamp
```

FINISHED



settings ▲

All fields:

status avgMeasuredTime avgSpeed extID medianMeasuredTime TIMESTAMP vehicleCount _id

REPORT_ID

Keys

status ✕

Groups

Values

vehicleCount SUM ✕



Results are limited by 1000.

Took 1 min 43 sec. Last updated by anonymous at March 31 2017, 1:15:18 PM. (outdated)

```
1 %pyspark
2
3 avgSpeed_corr = lambda x: x.corrwith(x['avgSpeed'])
4 by_hour = roadtraffic.groupby(lambda x: x.hour)
5 by_year.apply(avgSpeed_corr)
6 import statsmodels.api as sm
7 def regression(data, yvar, xvars):
8     Y = data[yvar]
9     X = data[xvars]
10    X['intercept'] = 1.
11    result = sm.OLS(Y,X).fit()
12    return result.params
13
14 by_year.apply(regression, 'AAPL', ['SPX'])
15
```

ERROR

Took 0 sec. Last updated by anonymous at March 31 2017, 1:17:39 PM. (outdated)

```
1 %pyspark
2 grouped = roadtraffic.groupby('hour')
3 get_wavg = lambda g: np.average(g['avgSpeed'], weights=g['vehicleCount'])
4 grouped.apply(get_wavg)
```

FINISHED

```
07 49.978347
00 51.107022
```

```
08 51.197955
09 50.778213
10 50.563207
11 50.482665
12 49.883909
13 47.909229
14 47.336310
15 50.606387
16 53.290080
17 54.765656
18 55.576947
19 55.817952
20 55.548410
21 56.290607
22 57.216225
23 58.958561
data = fit + 64
```

Took 12 sec. Last updated by anonymous at March 31 2017, 1:26:11 PM.

```
1 %pyspark
2 import statsmodels.api as sm
3
4 def regression(data, yvar, xvars):
5     Y = data[yvar]
6     X = data[xvars]
7     X['intercept'] = 1.
8     result = sm.OLS(Y,X).fit()
9     return result.params
10 by_hour.apply(regression,'avgSpeed',['vehicleCount'])
```

ERROR

Took 16 sec. Last updated by anonymous at March 31 2017, 1:25:27 PM.

```
%r
setwd("~/Downloads/traffic_feb_june")
files = list.files(getwd())
library(readr)
library(dplyr)
tbl = lapply(files, read_csv) %>% bind_rows()
summary(tbl)
Mode :character      Median : 80.0      Median : 43.00      Median : 825.0
                        Mean : 106.8      Mean : 43.94      Mean : 827.9
                        3rd Qu.: 116.0      3rd Qu.: 57.00      3rd Qu.: 938.0
                        Max. :3656.0      Max. :150.00      Max. :1058.0

medianMeasuredTime      TIMESTAMP                        vehicleCount
Min. : 0.0      Min. :2014-02-13 11:30:00      Min. : 0.000
1st Qu.: 53.0      1st Qu.:2014-03-30 06:30:00      1st Qu.: 0.000
Median : 80.0      Median :2014-05-11 09:55:00      Median : 1.000
Mean : 106.8      Mean :2014-05-30 14:03:59      Mean : 3.182
3rd Qu.: 116.0      3rd Qu.:2014-08-18 12:35:00      3rd Qu.: 4.000
Max. :3656.0      Max. :2014-09-30 23:55:00      Max. :121.000

_id      REPORT_ID
Min. : 189942      Min. :158324
1st Qu.: 5568083      1st Qu.:184621
Median :10944545      Median :190770
Mean :13157424      Mean :190461
3rd Qu.:22800528      3rd Qu.:197977
Max. :28061521      Max. :210100
```

FINISHED


MUX. : 20004021 MUX. : 210199

Took 1 min 31 sec. Last updated by anonymous at April 23 2017, 11:26:27 PM.

FINISHED

```
%r
tbl$hour<- as.POSIXlt(tbl$TIMESTAMP)$hour
tbl$Date<- as.Date(tbl$TIMESTAMP)

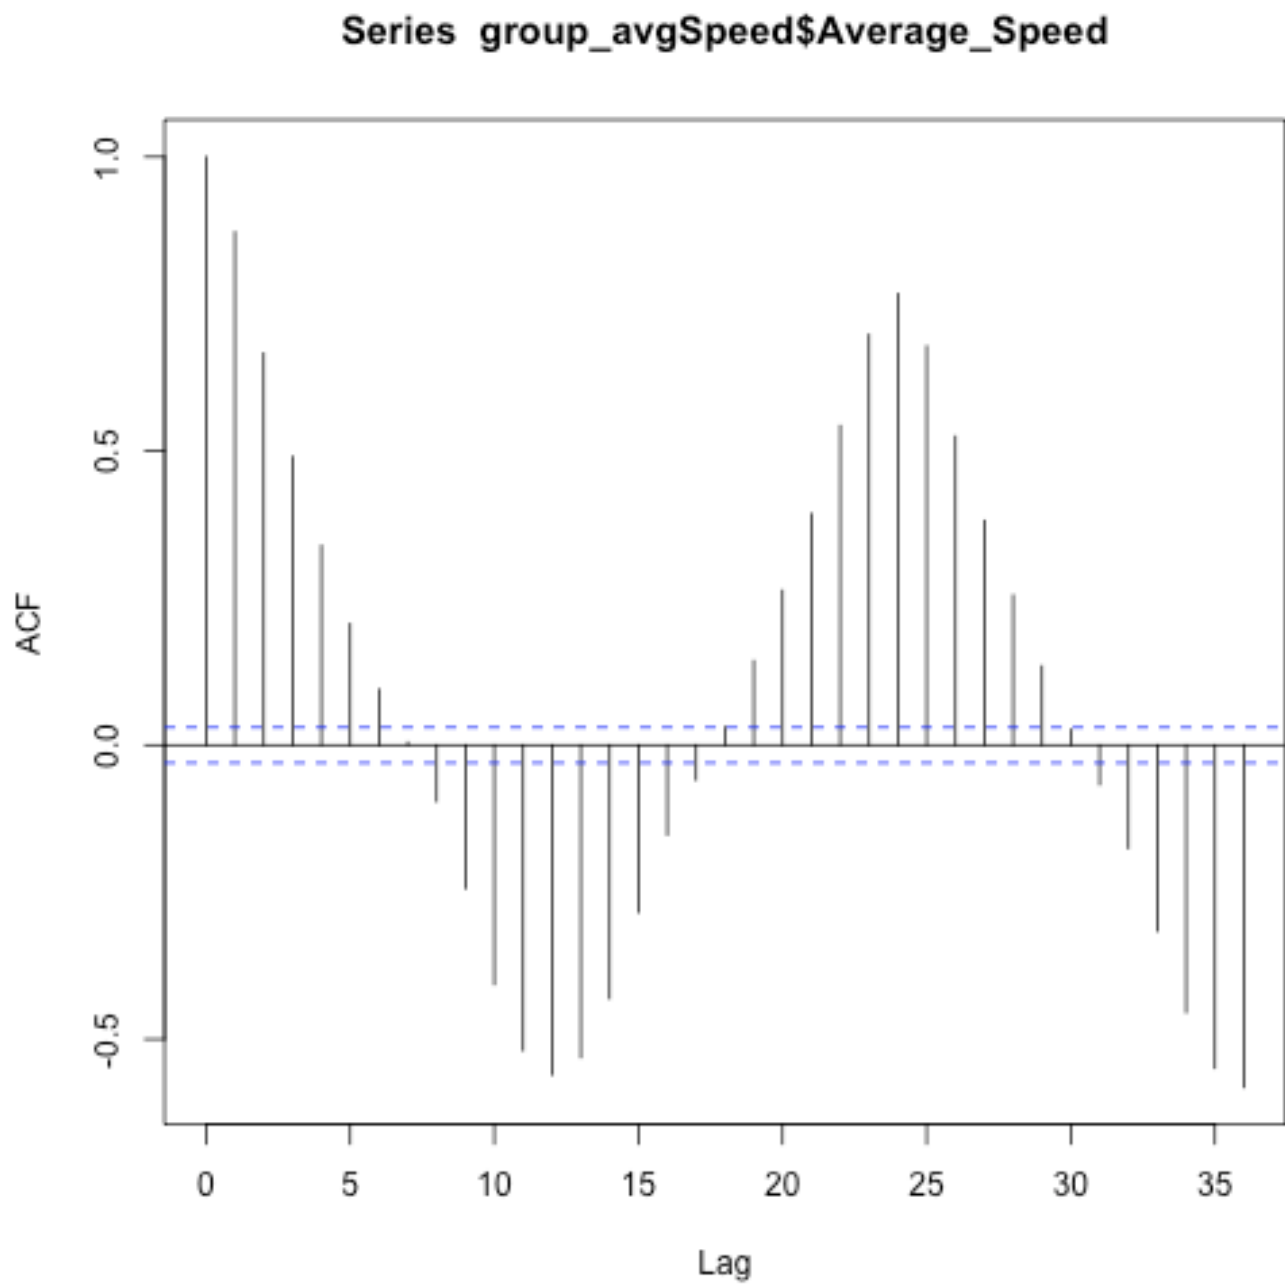
group_avgSpeed <- aggregate(tbl$avgSpeed, by = list(tbl$Date, tbl$hour), FUN = mean)
colnames(group_avgSpeed) <- c("Date", "Hour", "Average_Speed")
group_avgSpeed <- group_avgSpeed[with(group_avgSpeed, order(Date, Hour)), ]
group_avgSpeed$Date_Hour <- paste(as.character(group_avgSpeed$Date), paste(as.character(group.
plot(group_avgSpeed$Average_Speed, type = 'l')
```



Took 4 min 18 sec. Last updated by anonymous at April 23 2017, 11:31:51 PM.

```
%r
acf(group_avgSpeed$Average_Speed)
```

FINISHED




Took 0 sec. Last updated by anonymous at April 23 2017, 11:32:37 PM.

```
%r
pacf(group_avgSpeed$Average_Speed)
```

FINISHED

Series group_avgSpeed\$Average_Speed




Took 1 sec. Last updated by anonymous at April 23 2017, 11:32:52 PM.


```
%r
library(forecast)
train_avgSpeed <- group_avgSpeed[-c(4187:4198),]
fit_avgSpeed <- arima(train_avgSpeed$Average_Speed,order = c(4,2,7))
tsdiag(fit_avgSpeed)
```

FINISHED


Standardized Residuals



ACF of Residuals



p values for Ljung-Box statistic



Took 4 sec. Last updated by anonymous at April 23 2017, 11:36:04 PM.

FINISHED

```
%r
summary(fit_avgSpeed)

Call:
arima(x = train_avgSpeed$Average_Speed, order = c(4, 2, 7))
Coefficients:
      ar1      ar2      ar3      ar4      ma1      ma2      ma3      ma4
  1.1006 -1.0084  0.3534  0.1155 -1.7851  1.2711 -0.4577 -0.4285
s.e.  0.0891  0.1336  0.1305  0.0682  0.0878  0.1929  0.1847  0.0964
      ma5      ma6      ma7
  0.2477  0.0539  0.1061
s.e.  0.0459  0.0522  0.0218
sigma2 estimated as 2.471:  log likelihood = -7833.14,  aic = 15690.29
Training set error measures:
```


	ME	RMSE	MAE	MPE	MAPE	MASE
Training set	-0.005745859	1.571492	1.082008	-0.1138453	2.518093	0.9114698

ACF1

Training set -0.005423384

Took 0 sec. Last updated by anonymous at April 23 2017, 11:36:08 PM.

```
%r
pred_speed<- predict(fit_avgSpeed, n.ahead=12)
TS_result<-cbind(pred_speed$pred[1:12], group_avgSpeed$Average_Speed[c(4187:4198)])
colnames(TS_result) <- c("Predicted", "Actual")
TS_result
TS_result<-as.data.frame(TS_result)
plot(TS_result$Actual,type='l', col = 'red', ylim=c(25,50), xlab = 'Time', ylab = 'Average Speed')
lines(TS_result$Predicted,type='l', col = 'blue')
legend('topright', names(TS_result) , lty=1, col=c('blue', 'red'), bty='n', cex=.75)
```

FINISHED

Predicted	Actual
-----------	--------

[1,]	38.79589	38.83837	[2,]	37.89249	36.03330	[3,]	37.53761	34.39095	[4,]	38.60748	37.37238	[5,]	40.15553	
41.74689	[6,]	40.58624	42.28561	[7,]	39.73291	43.07860	[8,]	38.92659	43.89440	[9,]	39.12721	44.53928	[10,]	39.80582
45.30594	[11,]	39.86345	45.46067	[12,]	39.11689	45.83157								